Student ID: 406410035
Name: 秦紫頤
E-mail: chinjoyce30@gmail.com

Lab Title: Exception

Lab Purpose:

First, observe the initialization of the exception. Second, observe the implementation of system calls. Third, add a new system call.

Lab Procedure:

1. Observe exception handle
    a. Download Linux kernel source code for Raspberry Pi
    b. Add a printk statement **“Initialize traps\n”** into ~/pi_kernel/linux/init/main.c
    c. Add a printk statement **“arm: system call handler initialization -> see assembly code\n”** in the function void __init trap_init(void)
    d. Compile the kernel:
        i. cd ~/pi_kernel/linux
        ii. make ARCH=arm bcm2709_defconfig
        iii. make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- bzImage
    e. Copy the compiled zImage to SD card
        i. Plug the SD card into your laptop
        ii. mount /dev/sdb1 ~/mmc1
        iii. sudo rm ~/mmc1/zImage
        iv. cd ~/pi/kernel/linux/arch/arm/boot
        v. sudo cp -rf zImage ~/mmc1
        vi. sudo umount ~/mmc1
    f. Put the SD card back to Pi and power on
    g. dmesg

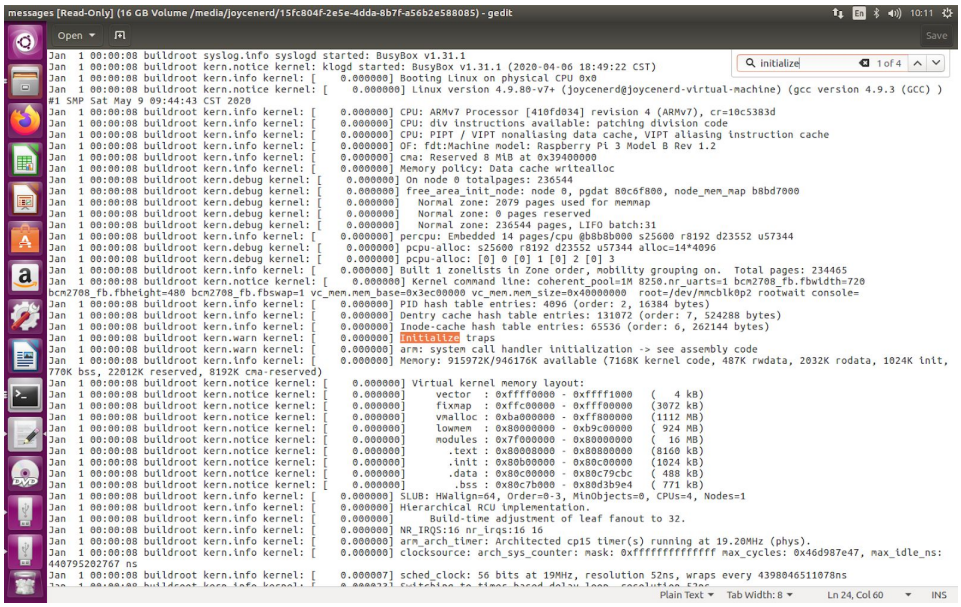

2. Observe the system call mechanism
    a. Write a simple hello world program using C
    b. Use cross compiler to compile your program: arm-linux-gnueabihf-gcc -static hello.c -o hello.exe
    c. Disassemble the program: arm-linux-gnueabihf-objdump -d hello.exe > assembly

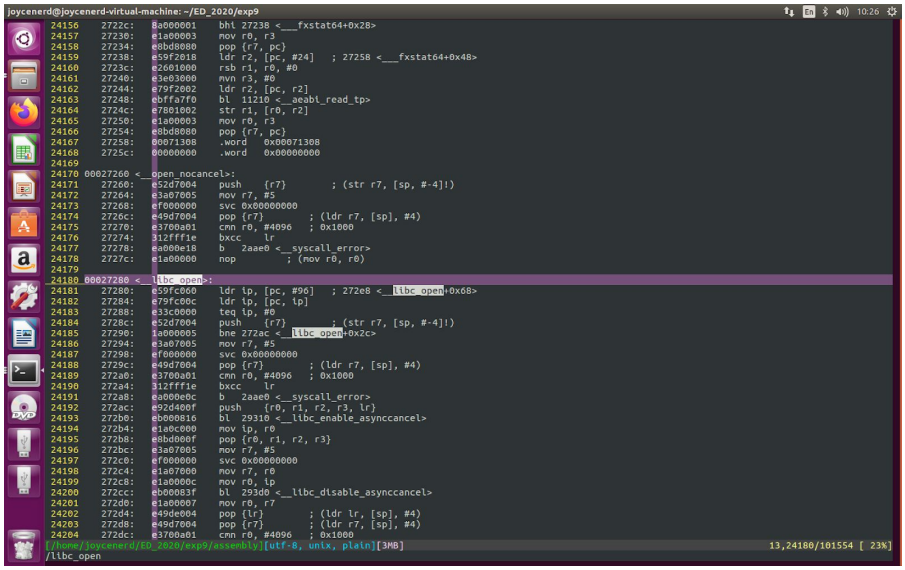

3. Write your own system call
    a. Add **“CALL(sys_mysyscall)”** to ~/pi_kernel/linux/arch/arm/kernel/calls.S
    b. Add define statement of the system call in unistd.h: **#define __NR_mysyscall (__NR_SYSCALL_BASE+397)**
    c. Write the content of the new added system call into ~/pi-kernel/linux/arch/arm/kernel/mysyscall.c
    d. Add the function declaration in syscalls.h: **asmlinkage void sys_mysyscall(int a, char* b)**;
    e. Modified the Makefile in ~/pi_kernel/linux/arch/arm/kernel: add **“my_syscall.o”** at the tail of **obj-y**
    f. Compile the kernel again and copy the zImage to the SD card
4. Test the system call
    a. Write a user program (mytestsys.c) to test call the system call I added in the previous part
    b. Use cross compiler to compile the program arm-linux-gnueabihf-gcc -I ~/pi-kernel/linux/include/ -static -g mytestsys.c -o mytestsys.exe
    c. Copy mytestsys.exe into the second partition of the SD card
    d. Put the SD card back into Pi and power on
    e. Execute mytestsys.exe to see the system call we added has worked or not

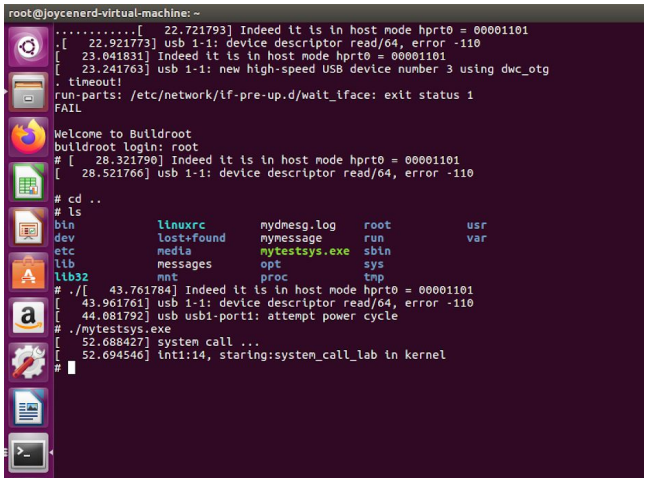

Problems and Discussions

- Q&A
  What is asmlinkage?
  This tells the compiler that the function should not expect to find any of its arguments in registers, but only on the CPU's stack. All system calls are marked with the asmlinkage tag, so they all look to the stack for the arguments.

- Discussions

  The first problem I encountered is that when using extra monitor and keyboard to display the message of Pi, there is an error about the device, which **doesn't allow the keyboard and mouse or any other extra device**. This error affects a lot cause this means we can't log in to the system in Pi. Another classmate found out that when **using USB to TTL serial transmission cable and remote login the system** although it still shows the error, in this way I can log in to the system cause I'm not plugging an extra device into Pi.

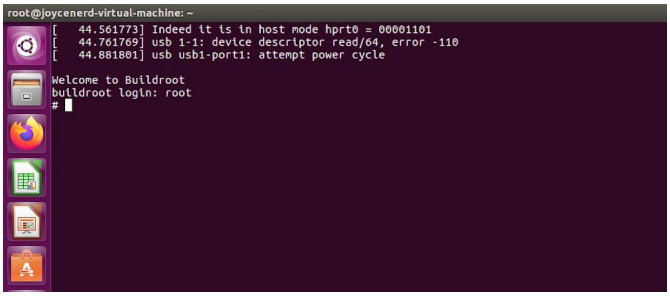

The second problem is after typing dmesg to see if "Initialize trap" and "arm: system call handler initialization -> see assembly code\n" are shown on the screen. There are so many messages and I can't

roll up to see the previous page. Then I find where the dmesg is saved in, which contains the original dmesg in it. It is in **/var/log/message**s. I search for the specific message inside that file and it actually works.